

GRAMPA: An Esoteric Programming Language to Simulate Parallel Computing

Contact: [email]@grinnell.edu

Reilly Grant
Grinnell College
[grantrei]

Tristan Knoth
Grinnell College
[knothtri17]

Chris Kottke
Grinnell College
[kottkech17]

Abstract

Parallel computing is an important way to process over very large problems. However, developing parallel applications is extremely difficult. We develop Generalized Rick And Morty ProgrAMming (GRAMPA), an esoteric imperative programming language supporting a simple forking model in order to simulate parallelism and introduce students shared memory and other rudimentary parallel computing concepts. The language supports only a limited set of commands, and includes syntax is based on the popular cartoon Rick and Morty in order to present parallel computing concepts in a fun and accessible manner.

1. Introduction

As parallelism becomes the most important paradigm for large-scale information processing, developing parallel applications is becoming an increasingly important skill for any developer. Students who begin to think about splitting problems up between processors earlier in their computer science education will potentially see more success learning more sophisticated parallelism paradigms later on. To this end, the popular TV show Rick and Morty presents the perfect medium through which to introduce students to rudimentary parallelism concepts. We develop GRAMPA, a simple Turing Complete imperative language with syntax based on references to Rick and Morty that supports a simple model of forking across shared memory. In Rick and Morty, the main characters travel between dimensions. The show's clear conceptual connections to multithreading may help alleviate the pain of learning to parallelize simple algorithms. GRAMPA's syntax is intended to look something like written natural language.

2. Prior Work

GRAMPA relies on a number of existing technologies, particularly the Haskell Parsec library. Our parser is built in Haskell using Parsec, which allows users to combine parsers via monads.... Didn't really know what else to put in this section.

3. Examples

4. Parsing GRAMPA

Using Parsec, we can build a recursive descent parser for GRAMPA by combining parsers for different types of expressions and statements in our language. Given a grammar for our language, we can define parsers for each variable in the grammar and thus recursively parse the entire language.

Grammar

The following Context-Free Grammar defines the GRAMPA syntax, and indicates the parsing hierarchy. The variable *STRING* refers to any string consisting only of chars, and *INT* refers to any integer.

$$\begin{aligned} S &\rightarrow UNIV \\ UNIV &\rightarrow \text{universe } STRING \text{ } STMT \text{ destroy universe } | UNIV UNIV \\ STMT &\rightarrow PORTAL | IF | DECL | PRINT | WHILE | STMT ST \\ EXPR &\rightarrow OP1 \\ OP1 &\rightarrow AND | OR | OP2 \\ OP2 &\rightarrow NUMEQ | NUMLT | NUMGT | OP3 \\ OP3 &\rightarrow ADD | SUB | OP4 \\ OP4 &\rightarrow MUL | DIV | MOD | TERM \\ TERM &\rightarrow BASE | PARENA | PARENB | STRING \\ BASE &\rightarrow INT | BOOL \\ BOOL &\rightarrow \text{right} | \text{wrong} \\ PORTAL &\rightarrow \text{lets grab our } STRING \text{ and portal out of here} \\ IF &\rightarrow \text{if } OP1 \text{ then } STMT \text{ otherwise } STMT \text{ wubulubadubdub} \\ DECL &\rightarrow STRING \text{ means } EXPR \\ PRINT &\rightarrow \text{show me } STRING \\ WHILE &\rightarrow \text{while } OP1 \text{ do this for grandpa } STMT \text{ thanks Summer} \\ AND &\rightarrow OP2 \text{ and } OP1 \\ OR &\rightarrow OP2 \text{ or } OP1 \\ NUMEQ &\rightarrow OP3 \text{ is the same as } OP3 \\ NUMLT &\rightarrow OP3 \text{ is less than } OP3 \\ NUMGT &\rightarrow OP3 \text{ is greater than } OP3 \\ ADD &\rightarrow OP4 \text{ plus } OP3 \\ SUB &\rightarrow OP4 \text{ minus } OP3 \\ MUL &\rightarrow TERM \text{ times } OP4 \\ DIV &\rightarrow TERM \text{ divided by } OP4 \end{aligned}$$

$MOD \rightarrow TERM \text{ mod } OP4$
 $PARENA \rightarrow \text{you gotta } OP3 \text{ Morty}$
 $PARENB \rightarrow \text{you gotta } OP1 \text{ Morty}$

Given the context-free grammar above, we define parsers for each of the individual substitution rules. For example, to parse a multiplication, we look for a *TERM* on the left side of the expression, a "times" to indicate that we are multiplying two expressions, and an *OP4* on the right hand side of the expression. In Haskell, this is implemented as follows:

Algorithm 1: Multiplication Parser

```

whitespace
e1 ← termParser
whitespace
string "times"
whitespace
e2 ← op4Parser
whitespace
return $ EBin Mul e1 e2

```

The sequence of instructions above is wrapped in a "do" block to create a parser for multiplication. *whitespace* is a parser designed to consume all whitespace. As such, our language is completely whitespace insensitive as long as one separates commands by any amount of whitespace. *string* is a parser built into Parsec, which parses a specific string. In this way, we combine parsers for different types of expressions and recursively parse the entire document.

Parsec also gives the user the option to "try" a parser. If one uses "try" when calling a specific parser, the parser will only consume input if it successfully parses the entire expression. Thus, if it encounters an error it will attempt again to parse the exact same string for a different type of expression. This feature allows us to combine parsers for different types of expressions quite easily. For example, we parse statements, which have a number of forms, in the following way:

```

stmt = try sPortal <|> try sIf <|> try sDec <|>
try sPrint <|> try sWhile

```

Thus, if the parser does not find an "if" statement, it can move on to look for a variable declaration or a print statement, for example. "Trying" different parsers in this way allows us to combine parsers like the multiplication parser described above and recursively parse the entire context-free grammar.

As a final note, one should observe that our order of operations is implicitly embedded in the grammar, and therefore the parser. The relevant portion of the grammar begins with the $EXPR \rightarrow OP1$ substitution. Here, we parse operations in ascending precedence. We parse the terminals and parenthesized expressions last in order to ensure that they are leafs in the parse tree and are therefore evaluated first.

4.1 Left Associativity

Parsing a file using the process described above results in a right-associative parse tree. This means that the expression $1 + 2 + 3$ is parsed as $1 + (2 + 3)$ rather than $(1 + 2) + 3$, which is what one would expect. In many cases this difference is nontrivial. For example, we need to ensure that $16/4/2$ evaluates to $(16/4)/2 = 2$ rather than $16/(4/2) = 8$. To rectify this problem, we transform the abstract syntax tree generated by our parser to reorder operations and guarantee left-associativity.

5. GRAMPA Code Generation and Execution

6. Parallelism in GRAMPA

Parallelism is an important concept and often is difficult to grasp for people who are new to it. People who are used to standard functional and imperative styles are often tripped up by the indeterminacy that parallelism introduces. In order to aid users who are unaccustomed to parallelism and its pitfalls and to cohere with the style of GRAMPA, we have implemented a mechanism for simulated parallelism which allows users to run programs on an indeterminate number of parallel threads. By simulating threads rather than actually implementing multithreading, we allow the user to explore this concept in a relatively consequence free environment. Furthermore, by controlling how the simulated multithreading works, we can present a simplistic representation of the concept that does not require the user to fight with more advanced topics such as context switching, locks, and signals.

Each parallel thread is represented by a different universe where the first universe listed is initially run in a single threaded environment. Universes are defined as follows:

```
universe name [STMT] destroy universe
```

In order to initiate multithreaded executions of multiple universes and essentially fork the program execution, the following Portal invocation is used:

```
lets grab our universe name and portal out of here
```

Upon executing this invocation, execution of the new universe will follow in parallel with the current executions. Rather than implementing true parallelism however, a mechanism similar to forking in C is used. In order to run this, the statement list of each universe currently being executed is stored in a list. In each step, the head statement of the head universe is popped and executed. Following this, the head universe is moved to the tail of the list. In this way, universes are cycled through and one universe will execute one instruction after all others have also executed one instruction. This workflow is changed slightly when the current executed instruction is a Portal instruction. In this case, the portal instruction returns the statement list of the new universe and this is appended to the end of the universe list followed by the universe which called the Portal instruction. When a universe has no more instructions, it is removed from the universe list. When there are no more universes in the list, execution ends.

Importantly, because all variables declared have multiversal scope, different universes can interact with and alter variables declared and used in other universes. This creates the potential for race conditions and allows users to implement their own basic locks in order to avoid these race conditions. An example of this can be found in code.txt. In addition to finding all prime numbers less than 1,000, this program also finds the sum of all of these primes. We can do this by declaring a variable *sum* and adding each prime we find to it. Note however that we cannot simply print *sum* once universe one finishes as universe two may still be processing; rather we must wait until all universes have terminated. To implement this, we introduce the variable *uTwoNotDone* which is initialized to right. At the end of universe one, we create a while loop which constantly checks the value of *uTwoNotDone*. If *uTwoNotDone* is still right, we perform an action which amounts to doing nothing and check again. At the end of universe two, we set *uTwoNotDone* to wrong in order to mark that universe two has finished processing. At this point, universe one escapes its loop and prints the value of *sum*.

7. Conclusion

A. Appendix Title

This is the text of the appendix, if you need one.

Acknowledgments

Acknowledgments, if needed.

References

P. Q. Smith, and X. Y. Jones. ...reference text...